

Practical 1 – Simple Linear Regression

- `import numpy as np`
`import matplotlib.pyplot as plt`
`import pandas as pd`
`dataset = pd.read_csv('Salary_Data.csv')`
`dataset.head()`
- `X = dataset.iloc[:, :-1].values`
`y = dataset.iloc[:, -1].values`
- `print(X)`
- `print(y)`
- `from sklearn.model_selection import train_test_split`
`X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 1/3, random_state=0)`
- `print(X_train)`
- `print(X_test)`
- `print(y_train)`
- `print(y_test)`
- `from sklearn.linear_model import LinearRegression`
`regressor = LinearRegression()`
`regressor.fit(X_train, y_train)`
- `y_pred = regressor.predict(X_test)`
`y_pred`
- `y_test`
- `plt.scatter(X_train, y_train, color = 'red')`
`plt.plot(X_train, regressor.predict(X_train), color = 'blue')`
`plt.title('Salary vs Experience (Training set)')`
`plt.xlabel('Years of Experience')`
`plt.ylabel('Salary')`
`plt.show()`
- `plt.scatter(X_test, y_test, color = 'red')`
`plt.plot(X_test, regressor.predict(X_test), color = 'blue')`

```
plt.title('Salary vs Experience (Test set)')
plt.xlabel('Years of Experience')
plt.ylabel('Salary')
plt.show()
```

Practical 2 – Multiple Linear Regression

- `import numpy as np`
`import matplotlib.pyplot as plt`
`import pandas as pd`
`dataset = pd.read_csv('50_Startups-2.csv')`
`dataset.head()`
- `X = dataset.iloc[:, :-1].values`
`y = dataset.iloc[:, -1].values`
- `print(X)`
- `print(y)`
- `from sklearn.compose import ColumnTransformer`
`from sklearn.preprocessing import OneHotEncoder`
`ct = ColumnTransformer(transformers = [('encoder', OneHotEncoder(), [3])], remainder = 'passthrough')`
`X = np.array(ct.fit_transform(X))`
- `print(X)`
- `from sklearn.model_selection import train_test_split`
`X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 0)`
- `from sklearn.linear_model import LinearRegression`
`regressor = LinearRegression()`
`regressor.fit(X_train, y_train)`
- `y_pred = regressor.predict(X_test)`
`np.set_printoptions(precision = 2)`
`print(np.concatenate((y_pred.reshape(len(y_pred), 1), y_test.reshape(len(y_test), 1)), 1))`

Practical 3 – Logistic Regression

- `import pandas as pd`
`iris_data = pd.read_csv('Iris.csv')`
`iris_data.head()`
- `iris_data.info()`
- `from sklearn.preprocessing import LabelEncoder`
`encoder = LabelEncoder()`
`iris_data['Species'] = encoder.fit_transform(iris_data['Species'])`
- `iris_data.head(150)`
- `import matplotlib.pyplot as plt`
`plt.pie(iris_data['Species'].value_counts(),labels=['Setosa','Versicol
or','Virginica'],autopct='%0.2f')`
`plt.show()`
- `x = iris_data.drop('Species',axis=1)`
`y=iris_data['Species']`
- `print(x)`
- `print(y)`
- `from sklearn.model_selection import train_test_split`
`x_train,x_test,y_train,y_test =`
`train_test_split(x,y,test_size=0.2,random_state=2)`
- `from sklearn.linear_model import LogisticRegression`
`model = LogisticRegression(max_iter=1000)`
`model.fit(x_train,y_train)`
- `pred_train = model.predict(x_train)`
- `from sklearn.metrics import confusion_matrix,accuracy_score`
`accuracy_score(y_train,pred_train)`
- `pred_test = model.predict(x_test)`
- `accuracy_score(y_test,pred_test)`
- `confusion_matrix(y_test,pred_test)`

Practical 4 – KNN

- `import numpy as np`
`import pandas as pd`
`import matplotlib.pyplot as plt`
- `dataset = pd.read_csv('Social_Network_Ads.csv')`
`X = dataset.iloc[:, :-1].values`
`y = dataset.iloc[:, -1].values`
- `from sklearn.model_selection import train_test_split`
`X_train, X_test, y_train, y_test =`
`train_test_split(X, y, test_size=0.25, random_state=0)`
- `print(X_train)`
- `print(X_test)`
- `print(y_train)`
- `print(y_test)`
- `from sklearn.preprocessing import StandardScaler`
`sc = StandardScaler()`
`X_train = sc.fit_transform(X_train)`
`X_test = sc.transform(X_test)`
- `print(X_train)`
- `print(X_test)`
- `from sklearn.neighbors import KNeighborsClassifier`
`classifier = KNeighborsClassifier(n_neighbors = 5, metric`
`= 'minkowski', p = 2)`
`classifier.fit(X_train, y_train)`
- `print(classifier.predict(sc.transform([[40, 200000]])))`
- `y_pred = classifier.predict(X_test)`
`print`
`(np.concatenate((y_pred.reshape(len(y_pred), 1), y_test.reshape(len`
`(y_test), 1)), 1))`
- `from sklearn.metrics import confusion_matrix, accuracy_score`
`cm = confusion_matrix(y_pred, y_test)`
`print(cm)`
`accuracy_score(y_pred, y_test)`

Practical 5 – SVM

- `import numpy as np`
`import pandas as pd`
`import matplotlib.pyplot as plt`
- `dataset = pd.read_csv('Social_Network_Ads.csv')`
`X = dataset.iloc[:, :-1].values`
`y = dataset.iloc[:, -1].values`
- `from sklearn.model_selection import train_test_split`
`X_train, X_test, y_train, y_test =`
`train_test_split(X, y, test_size=0.25, random_state=0)`
- `print(X_train)`
- `print(X_test)`
- `print(y_train)`
- `print(y_test)`
- `from sklearn.preprocessing import StandardScaler`
`sc = StandardScaler()`
`X_train = sc.fit_transform(X_train)`
`X_test = sc.transform(X_test)`
- `print(X_train)`
- `print(X_test)`
- `from sklearn.svm import SVC`
`classifier = SVC(kernel='linear')`
`classifier.fit(X_train, y_train)`
- `print(classifier.predict(sc.transform([[40, 200000]])))`
- `y_pred = classifier.predict(X_test)`
`print`
`(np.concatenate((y_pred.reshape(len(y_pred), 1), y_test.reshape(len`
`(y_test), 1)), 1))`
- `from sklearn.metrics import confusion_matrix, accuracy_score`
`cm = confusion_matrix(y_pred, y_test)`
`print(cm)`
`accuracy_score(y_pred, y_test)`

Practical 6 – K-Means

- `import numpy as np`
`import matplotlib.pyplot as plt`
`import pandas as pd`
- `dataset = pd.read_csv('Mall_Customers.csv')`
`X = dataset.iloc[:,[3,4]].values`
`print(X)`
- `from sklearn.cluster import KMeans`
`kmeans = KMeans(n_clusters = 5,init = 'k-means++',random_state=42)`
`y_kmeans = kmeans.fit_predict(X)`
`print(y_kmeans)`
- `plt.scatter(X[y_kmeans == 0,0],X[y_kmeans == 0,1], s = 100 , c = 'red', label='Cluster 1')`
`plt.scatter(X[y_kmeans == 1,0],X[y_kmeans == 1,1], s = 100 , c = 'blue', label='Cluster 2')`
`plt.scatter(X[y_kmeans == 2,0],X[y_kmeans == 2,1], s = 100 , c = 'green', label='Cluster 3')`
`plt.scatter(X[y_kmeans == 3,0],X[y_kmeans == 3,1], s = 100 , c = 'cyan', label='Cluster 4')`
`plt.scatter(X[y_kmeans == 4,0],X[y_kmeans == 4,1], s = 100 , c = 'magenta', label='Cluster 5')`
`plt.scatter(kmeans.cluster_centers_[0,0],kmeans.cluster_centers_[0,1],s=300,c='yellow',label='Centroids')`
`plt.title('Clusters of customers')`
`plt.xlabel('Annual Income(k$)')`
`plt.ylabel('Spending Score (1-100)')`
`plt.legend()`
`plt.show()`

Practical 7 – Hierarchical Clustering

- `import numpy as np`
`import matplotlib.pyplot as plt`
`import pandas as pd`
- `dataset = pd.read_csv('Mall_Customers.csv')`
`X = dataset.iloc[:,[3,4]].values`
`print(X)`
- `from sklearn.cluster import AgglomerativeClustering`
`Agg_hc = AgglomerativeClustering(n_clusters = 5, affinity = 'euclidean', linkage = 'ward')`
`y_hc = Agg_hc.fit_predict(X)`
`print(y_hc)`
- `plt.scatter(X[y_hc == 0, 0], X[y_hc == 0, 1], s = 100, c = 'red', label = 'Cluster 1') # plotting cluster 2`
`plt.scatter(X[y_hc == 1, 0], X[y_hc == 1, 1], s = 100, c = 'blue', label = 'Cluster 2') # plotting cluster 3`
`plt.scatter(X[y_hc == 2, 0], X[y_hc == 2, 1], s = 100, c = 'green', label = 'Cluster 3') # plotting cluster 4`
`plt.scatter(X[y_hc == 3, 0], X[y_hc == 3, 1], s = 100, c = 'cyan', label = 'Cluster 4') # plotting cluster 5`
`plt.scatter(X[y_hc == 4, 0], X[y_hc == 4, 1], s = 100, c = 'magenta', label = 'Cluster 5')`
`plt.title('Clusters of customers')`
`plt.xlabel('Annual Income (k$)')`
`plt.ylabel('Spending Score (1-100)')`
`plt.legend()`
`plt.show()`

Practical 8 – ANN

- `import numpy as np`
`import pandas as pd`
`import tensorflow as tf`

- `tf.__version__`
- `dataset = pd.read_csv('Churn_Modelling.csv')`
`X = dataset.iloc[:,3:-1].values`
`y = dataset.iloc[:,-1].values`
- `print(X)`
- `print(y)`
- `from sklearn.preprocessing import LabelEncoder`
`le = LabelEncoder()`
`X[:,2] = le.fit_transform(X[:,2])`
- `print(X)`
- `from sklearn.compose import ColumnTransformer`
`from sklearn.preprocessing import OneHotEncoder`
`ct =`
`ColumnTransformer(transformers=[('encoder',OneHotEncoder(),[1]`
`)],remainder = 'passthrough')`
`X = np.array(ct.fit_transform(X))`
- `print(X)`
- `from sklearn.model_selection import train_test_split`
`X_train,X_test,y_train,y_test=`
`train_test_split(X,y,test_size=0.2,random_state=0)`
- `from sklearn.preprocessing import StandardScaler`
`sc= StandardScaler()`
`X_train = sc.fit_transform(X_train)`
`X_test = sc.transform(X_test)`
- `ann = tf.keras.models.Sequential()`
- `ann.add(tf.keras.layers.Dense(units=6,activation='relu'))`
- `ann.add(tf.keras.layers.Dense(units=6,activation='relu'))`
- `ann.add(tf.keras.layers.Dense(units=1,activation='sigmoid'))`
- `ann.compile(optimizer='adam',loss = 'binary_crossentropy',metrics`
`= ['accuracy'])`
- `ann.fit(X_train,y_train,batch_size =32,epochs=100)`

- `print(ann.predict(sc.transform([(1,0,0,600,1,40,3,60000,2,1,1,500,00)]))>0.5)`
- `y_pred = ann.predict(X_test)`
`y_pred = (y_pred > 0.5)`
`print(np.concatenate((y_pred.reshape(len(y_pred),1),y_test.reshape(len(y_test),1)),1))`
- `from sklearn.metrics import confusion_matrix,accuracy_score`
`cm = confusion_matrix(y_test,y_pred)`
`print(cm)`
`accuracy_score(y_test,y_pred)`

Practical 9 – CNN

- `import tensorflow as tf`
`from keras.preprocessing.image import ImageDataGenerator`
- `tf.__version__`
- `train_datagen = ImageDataGenerator(rescale=1./225,`
`shear_range=0.2,`
`zoom_range=0.2,`
`horizontal_flip= True)`
`training_set`
`=train_datagen.flow_from_directory('drive/MyDrive/small_dataset/training_set',`
`target_size=(64,64),`
`batch_size = 32,`
`class_mode = 'binary')`
- `test_datagen = ImageDataGenerator(rescale = 1./255)`
`test_set =`
`test_datagen.flow_from_directory('drive/MyDrive/small_dataset/test_set',`
`target_size = (64,64),`
`batch_size = 32,`

class_mode = 'binary')

- `cnn = tf.keras.models.Sequential()`
- `cnn.add(tf.keras.layers.Convolution2D(filters=32,kernel_size=3,activation='relu',input_shape=[64,64,3]))`
- `cnn.add(tf.keras.layers.MaxPool2D(pool_size=2,strides=2))`
- `cnn.add(tf.keras.layers.Convolution2D(filters=32,kernel_size=3,activation='relu'))`
`cnn.add(tf.keras.layers.MaxPool2D(pool_size=2,strides=2))`
- `cnn.add(tf.keras.layers.Flatten())`
- `cnn.add(tf.keras.layers.Dense(units=128,activation='relu'))`
- `cnn.add(tf.keras.layers.Dense(units=1,activation='sigmoid'))`
- `cnn.compile(optimizer='adam',loss='binary_crossentropy',metrics=['accuracy'])`
- `cnn.fit(x=training_set,validation_data=test_set,epochs=25)`
- `import numpy as np`
`from tensorflow.keras.preprocessing import image`
`test_image =`
`image.load_img('drive/MyDrive/small_dataset/single_prediction/cat_or_dog_1.jpg',target_size=(64,64))`
`test_image = image.img_to_array(test_image)`
`test_image = np.expand_dims(test_image,axis=0)`
`result = cnn.predict(test_image)`
`training_set.class_indices`
`if result[0][0]==1:`
 `prediction = 'dog'`
`else:`
 `prediction = 'cat'`
- `print(prediction)`